


Token Economics





TOKEN
LEGEND


GitHub Copilot Runtime **Ported to Rust**

A faster, leaner foundation for agents



<https://aka.ms/copilot-migration>

Copilot runtime performance



■ Reduction ■ Remaining

91%

less memory

1,383 MB before / 126 MB after

65%

less processor time


312 s before / ~110 s after

4×

faster session lifecycle

5.25 s before / 1.33 s after

Cost savings



\$120K

total spend

\$0.14

per line of production Rust

136.3B

total tokens

130.6B
cached reads

4.2B
cache writes

0.9B
fresh input

0.6B
output

📈 **96.22% prompt-cache hit rate** — primary savings lever

🔗 128 PRs



= 4 tokens

```
{
  "BillingCurrency": "USD",
  "CustomerEntityId": "Default",
  "CustomerEntityType": "Retail",
  "Items": [
    {
      "currencyCode": "USD",
      "tierMinimumUnits": 0,
      "retailPrice": 1.2,
      "unitPrice": 1.2,
      "armRegionName": "eastus",
      "location": "US East",
      "effectiveStartDate": "2026-08-01T00:00:00Z",
      "meterId": "c6713056-7cef-508c-9e2e-7035984a4f44",
      "meterName": "5.6 luna ShortCo Opt Std G1 1M Tokens",
      "productId": "DZH318Z0T9WD",
      "skuId": "DZH318Z0T9WD/1166",
      "productName": "Azure OpenAI GPT5",
      "skuName": "5.6 luna ShortCo Opt Std G1",
      "serviceName": "Foundry Models",
      "serviceId": "DZH31600VF1F",
      "serviceFamily": "AI + Machine Learning",
      "unitOfMeasure": "1M",
      "type": "Consumption",
      "isPrimaryMeterRegion": true,
      "armSkuName": "56 luna ShortCo Opt Std G1"
    }
  ],
  "NextPageLink": null,
  "Count": 1
}
```

Currency: billing currency.

Retail price: list price charged.

Unit price: price per billing unit.

Region code: Azure region identifier.

Region name: human-friendly region name.

Effective date: when this price became active.

Meter name: full billing meter description.


Product name: product family.

SKU name: specific model/tier variant.

Service name: Azure service category.

Unit of measure: billed per 1M tokens.

Type: consumption-based billing.



<https://aka.ms/costs>

MODEL



Router



Triage

PROMPT



Step-back
planning

CONTEXT



Compression



RAG



Tool use



Caching

OUTPUT



Batching

Models



GPT-5.6 Luna

Small / Fast




GPT-5.6 Terra

Medium




GPT-5.6 Sol

Large / Reasoning



Triage



MODEL



Router

Triage

PROMPT



Step-back
planning

CONTEXT



Compression

RAG

Tool use


Caching

OUTPUT



Batching

Step-back planning



MODEL



☒ Router

☒ Triage

PROMPT



☒ Step-back
planning

CONTEXT



☒ Compression

☒ RAG

☒ Tool use


☒ Caching


OUTPUT




☒ Batching

Compression







Retrieval-Augmented Generation (RAG)



Tools



Caching



MODEL



☒ Router

☒ Triage

PROMPT



☒ Step-back
planning

CONTEXT



☒ Compression

☒ RAG

☒ Tool use


☒ Caching

OUTPUT



☒ Batching

Batching



MODEL



☒ Router

☒ Triage

PROMPT



☒ Step-back
planning

CONTEXT



☒ Compression

☒ RAG

☒ Tool use

☒ Caching

OUTPUT



☒ Batching

TETRIS + LUNA

Front-end screen guide

Real play. Real AI usage.

2. Score, lines, level

Current points, cleared rows and falling speed.

3. Hold

Save or swap a piece; once per locked piece.

4. Playfield

Place blocks to clear rows. Token labels follow pieces; the ghost marks a drop.

6. Game controls

Move, rotate, hold or drop during manual play.

7. Information & tools

Inspect prompts, costs and MCP; share or edit.

1. Session controls

Connection, Pause/Play and New game.

11. Room & allowance

Room code, online players, your sentence and remaining personal / shared AI tokens.

12. Leaderboard

Compare best points or points per cent of AI spend.

8. AI cost ticker

Est. total: \$0.03009738
183,386 tokens used.

9. Luna controls

Enable paid AI play; choose reasoning, MCP, compression and cache for the next call.

10. Request status


See request options, token counts and cache results.

5. Next

Preview the next three pieces.

Captured after a real Luna-assisted run: 6 lines cleared.

Luna is off here. Costs and savings are estimates, not an invoice.



AI COST TICKER + LUNA CONTROLS

A close-up of the real run: 183,386 AI tokens used.

AI cost ticker

AI COST / USD EST.
\$0.0300973
8
Live rates

LAST CALL
\$0.00150786
9,429 in / 32 out

AI TOKENS LEFT
816,614
183,386 used

SAVED / EST.
\$0.01761882
Compression + cache

Total AI cost: \$0.03009738
Reported input/output usage priced at current USD rates. Cumulative across games; not your Tetris score or an invoice.

Last call: \$0.00150786
Cost of the most recent completed Luna request. This reply used 9,429 input and 32 output tokens.

Saved (estimate): \$0.01761882
Compression savings plus the net effect of cache reads/writes. Writes can add cost; this is not a refund or bonus points.

AI tokens left: 816,614
The smaller of your personal and shared room balances. 183,386 input + output tokens used, including cached input.

Scoring is separate

Points / cent = best score / cumulative AI spend in US cents.
Zero or unknown AI spend is unranked, not infinite.

Luna controls

Each switch changes how the next request works.

🏠 **Ask Luna** ☐

MCP
Luna off ☐

🧠 **Reasoning** ☐

Reasoning tokens
(all calls)
0

Reasoning tokens
(last call, off)
0

📄 **Compression** ☐

Cache
Luna off ☐

Ask Luna: continuous paid play
The model keeps choosing and playing moves until stopped. Switch off, then press Play to resume manual control.

MCP: two-placement lookahead
A real tool forecasts the next two placements, including Hold. Adds analysis input tokens and latency; no survival guarantee.

Reasoning: optional low effort
Can increase output tokens and latency. Off uses no reasoning. The counters are already included in output and cost.

Cache: reuse fixed instructions
Provider-confirmed hits can be cheaper; writes may cost extra. Hits are not guaranteed. Cached tokens still use your allowance.

Compression: same board, fewer tokens
Losslessly packs the same board; no cells are removed. Affects future prompts, not tokens or costs already recorded.

Luna is off in this captured frame.

All four options are off and disabled until Ask Luna is on.
Saved preferences return. Earlier costs are not reset.
An in-flight request can still finish and be charged.

Magnified crops from the same real game. Displayed costs and savings are estimates.

TETRIS

SCORE
688

LINES
3

LEVEL
1

HOLD
_make

NEXT
_jmy

blocks
_blocks

10 x 20
LVL 01

10 x 20
LVL 01

Drop

Information & tools

MCP resultsCache activityCompressionAI costsLeaderboardShare gameChange sentence

AI COST / PENDING USAGE
\$0.0278802
Live rates
AI TOKENS LEFT
864 254
135 746 used

LAST CALL
\$0.0019422
11 427 in / 89 out
SAVED / EST.
\$0.004206
Compression + cache

Ask LunaReasoningCompressionMCPCache


Luna is thinking

Current request: Compression · Cache · Reasoning · MCP

-22% compression · 2 500 cached · 11 427 in / 89 out
Last call: \$0.0019422



<https://aka.ms/tokenfall>



■ Recorded Copilot cost

\$1,250.20


17.1% of chart total




■ Estimated cache savings

\$6,064.02


82.9% of chart total

Token Economics Checklist





<https://aka.ms/token-economics>



<https://github.com/roryp/token-design-patterns>

TETRIS

SCORE
688

LINES
3

LEVEL
1

HOLD

_world

01

10

_make

NEXT

_my

_blocks

10 x 20

LVL 01

...

Hello

_Tokens

.

!

Wow

↶

↷

↺

←

↓

→

↓ DROP

Information & tools

MCP results

Cache activity

Compression

AI costs

Leaderboard

Share game

Change sentence

AI COST / PENDING USAGE

\$0.0278802

Live rates

AI TOKENS LEFT

864 254

135 746 used

LAST CALL

\$0.0019422

11 427 in / 89 out

SAVED / EST.

\$0.004206

Compression + cache

Ask Luna

MCP

2-piece next

Reasoning

Low effort next

Reasoning tokens (all calls)

1 405

Reasoning tokens (last call, low)

47

Compression

Packed rows next

Cache

Reuse rules next

Luna is thinking

Current request: Compression · Cache · Reasoning · MCP

-22% compression · 2 500 cached · 11 427 in / 89 out

Last call: \$0.0019422



<https://github.com/roryp/tokenfall>