


Token Economics





**TOKEN
LEGEND**


GitHub Copilot Runtime **Ported to Rust**

A faster, leaner foundation for agents



<https://aka.ms/copilot-migration>

Copilot runtime performance



■ Reduction ■ Remaining

91%

less memory

1,383 MB before / 126 MB after

65%

less processor time


312 s before / ~110 s after

4×

faster session lifecycle

5.25 s before / 1.33 s after

Cost savings



\$120K

total spend

\$0.14

per line of production Rust

136.3B

total tokens

130.6B
cached reads

4.2B
cache writes

0.9B
fresh input

0.6B
output

📈 **96.22% prompt-cache hit rate** — primary savings lever

🔗 128 PRs



= 4 tokens

```
{
  "BillingCurrency": "USD",
  "CustomerEntityId": "Default",
  "CustomerEntityType": "Retail",
  "Items": [
    {
      "currencyCode": "USD",
      "tierMinimumUnits": 0,
      "retailPrice": 1.2,
      "unitPrice": 1.2,
      "armRegionName": "eastus",
      "location": "US East",
      "effectiveStartDate": "2026-08-01T00:00:00Z",
      "meterId": "c6713056-7cef-508c-9e2e-7035984a4f44",
      "meterName": "5.6 luna ShortCo Opt Std G1 1M Tokens",
      "productId": "DZH318Z0T9WD",
      "skuId": "DZH318Z0T9WD/1166",
      "productName": "Azure OpenAI GPT5",
      "skuName": "5.6 luna ShortCo Opt Std G1",
      "serviceName": "Foundry Models",
      "serviceId": "DZH31600VF1F",
      "serviceFamily": "AI + Machine Learning",
      "unitOfMeasure": "1M",
      "type": "Consumption",
      "isPrimaryMeterRegion": true,
      "armSkuName": "56 luna ShortCo Opt Std G1"
    }
  ],
  "NextPageLink": null,
  "Count": 1
}
```

Currency: billing currency.

Retail price: list price charged.

Unit price: price per billing unit.

Region code: Azure region identifier.

Region name: human-friendly region name.

Effective date: when this price became active.

Meter name: full billing meter description.

Product name: product family.


SKU name: specific model/tier variant.

Service name: Azure service category.

Unit of measure: billed per 1M tokens.

Type: consumption-based billing.





■ Recorded Copilot cost

\$1,250.20

17.1% of chart total

■ Estimated cache savings

\$6,064.02

82.9% of chart total

MODEL



☒ Router

☒ Triage

PROMPT



☒ Step-back
planning

CONTEXT



☒ Compression

☒ RAG


☒ Tool use

☒ Caching

OUTPUT







☒ Batching




<https://aka.ms/costs>

Token Economics Checklist





<https://aka.ms/token-economics>



<https://github.com/roryp/token-design-patterns>

TETRIS

SCORE
688

LINES
3

LEVEL
1

HOLD
_world

0110
_make

10
NEXT
_my

10
_blocks

10 x 20
LVL 01

⏮ ⏪ ⏩ ⏭ ⏮ ⏪ ⏩ ⏭

⏮ ⏪ ⏩ ⏭

Information & tools

MCP results

Cache activity

Compression

AI costs

Leaderboard

Share game

Change sentence

AI COST / PENDING USAGE

\$0.0278802

Live rates

AI TOKENS LEFT

864 254

135 746 used

LAST CALL

\$0.0019422

11 427 in / 89 out

SAVED / EST.

\$0.004206

Compression + cache

Ask Luna

MCP

2-piece next

Reasoning

Low effort next

Reasoning tokens (all calls)

1 405

Reasoning tokens (last call, low)

47

Compression

Packed rows next

Cache


Reuse rules next

Luna is thinking

Current request: Compression · Cache · Reasoning · MCP

-22% compression · 2 500 cached · 11 427 in / 89 out

Last call: \$0.0019422



<https://github.com/roryp/tokenfall>